


Model Architecture

I decided to go with the NVIDIA architecture as it is a model which is proven in practice. The architecture of the model is as follows.



Nvidia model - keras implementation.

```
model = Sequential()
    model.add(Lambda(lambda x: (x / 127.5) - 1., input_shape=(66, 200,
3)))
    model.add(Convolution2D(24, 5, 5, subsample=(2, 2), border_mode="valid", activation="elu"))
    model.add(Convolution2D(36, 5, 5, subsample=(2, 2), border_mode="valid", activation="elu"))
    model.add(Dropout(.5))
    model.add(Convolution2D(48, 5, 5, subsample=(2, 2), border_mode="valid", activation="elu"))
    model.add(Dropout(.5))
    model.add(Convolution2D(64, 3, 3, border_mode="valid", activation="elu"))
    model.add(Dropout(.5))
    model.add(Convolution2D(64, 3, 3, border_mode="valid", activation="elu"))
    model.add(Dropout(.5))
    model.add(Flatten())
    model.add(Dense(100,activation='elu'))
    model.add(Dropout(.25))
    model.add(Dense(50,activation='elu'))
    model.add(Dense(10,activation='elu'))
    model.add(Dense(1))
```

In [5]:

```
# Model details
# _____
# Layer (type)                Output Shape          Param #      Connected to
# =====
# lambda_1 (Lambda)           (None, 66, 200, 3)    0            lambda_input_1[0][0]
# _____
# convolution2d_1 (Convolution2D) (None, 31, 98, 24)    1824         lambda_1[0][0]
# _____
# convolution2d_2 (Convolution2D) (None, 14, 47, 36)    21636        convolution2d_1[0][0]
# _____
# dropout_1 (Dropout)          (None, 14, 47, 36)    0            convolution2d_2[0][0]
# _____
# convolution2d_3 (Convolution2D) (None, 5, 22, 48)     43248        dropout_1[0][0]
# _____
```

# dropout_2 (Dropout) 2d_3[0][0] #	(None, 5, 22, 48)	0	convolution
# convolution2d_4 (Convolution2D) [0][0] #	(None, 3, 20, 64)	27712	dropout_2
# dropout_3 (Dropout) 2d_4[0][0] #	(None, 3, 20, 64)	0	convolution
# convolution2d_5 (Convolution2D) [0][0] #	(None, 1, 18, 64)	36928	dropout_3
# dropout_4 (Dropout) 2d_5[0][0] #	(None, 1, 18, 64)	0	convolution
# flatten_1 (Flatten) [0][0] #	(None, 1152)	0	dropout_4
# dense_1 (Dense) [0][0] #	(None, 100)	115300	flatten_1
# dropout_5 (Dropout) [0] #	(None, 100)	0	dense_1[0]
# dense_2 (Dense) [0][0] #	(None, 50)	5050	dropout_5
# dense_3 (Dense) [0] #	(None, 10)	510	dense_2[0]
# dense_4 (Dense) [0] #	(None, 1)	11	dense_3[0]
=====			
# Total params: 252,219			
# Trainable params: 252,219			
# Non-trainable params: 0			

Attempts to reduce overfitting

Initially RELU activation were used for the architecture and using ELU provided better results on the track. Dropout is used to reduce overfitting of the model. Adam optimiser with a learning rate of .0001 was used.

Since the model is trained from scratch and the higher level features were not transferred using transfer learning, it became evident that desired model might not be the final model after the last epoch. The model was checkpointed after every epoch to examine the performance of the model on the track. It helps examine how the model is evolving over the epochs.

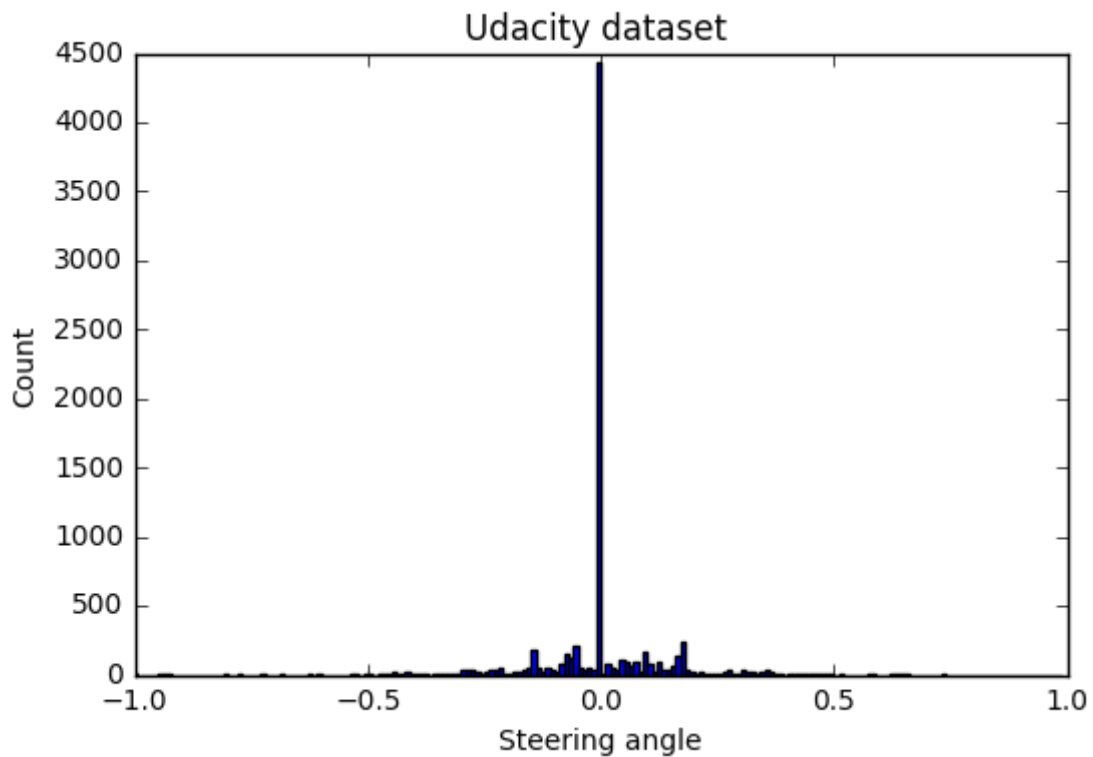
Model Parameters

- Samples per epoch - 81920
- Number of epoch - 10

Training Data Collection and Exploration

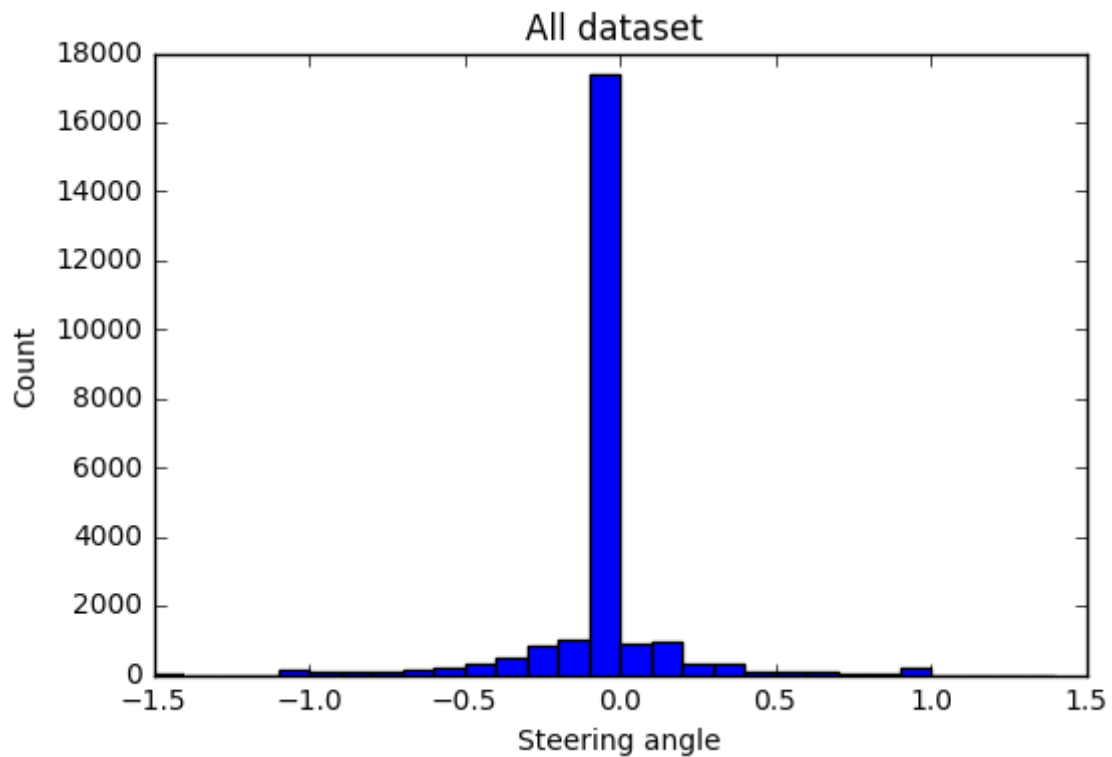
- The dataset provided by udacity is dominated by zero steering angles values.
 - When the model was trained using the udacity dataset. The car was not able to make sharp turns or recover back to the center of the road.
 - To augment the udacity dataset more general driving, sharp turns, and recovery video was collected.
 - The final dataset was an amalgamation of the above dataset.
-

Udacity provided data



Collected data + udacity data

Size of the dataset - 23771



Data Augmentation and preprocessing

- The image is normalised to have a range of $[-1 \ 1]$ using a keras lambda layer.
- YUV format is fed into the model.
- Random brightness augmentation
- Cropping the image
 - The images were cropped so the bonnet of the car, the sky and the tree are removed.
 - The images fed to the network contains only the road with a vantage point.
- A gaussian blur with a 3X3 kernel is applied to the image.
- one of the three cameras is chosen at random for the training set. when the left and right cameras are used a correction of 0.25 is used to correct the steering angle.
- while augmenting the steering angle the values are ensured to be between $[-1 \ 1]$
- Half the images in a batch are flipped horizontally and the angle multiplied by -1.0
- Images are shifted at random along the x axis.
- Data with throttle values below .25 is dropped from the dataset as it does not represent driving behaviour.

*****Only the training data is augmented. only unaugmented center camera is used for validation set.***

Augmented Images



Un augmented Images

